

Deep Scratch

Reusing Pretrained Layers

Transfer Learning

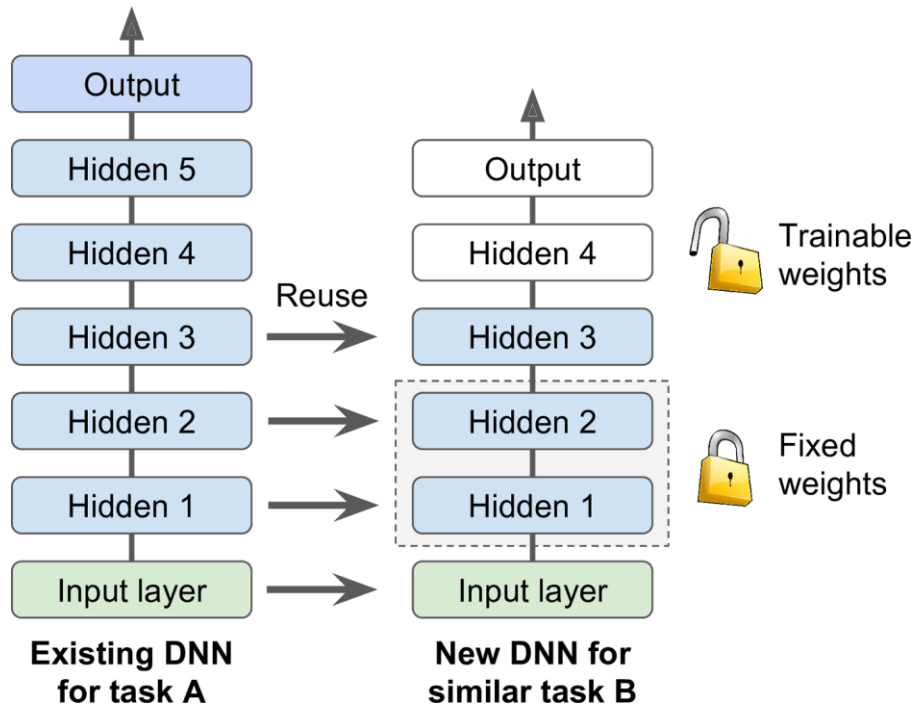
박 희 경

CONTENTS

1. Introduction
2. Freezing the Lower Layers
3. Tweaking, Dropping,
or Replacing the Upper Layers
4. Caching the Frozen Layers
5. Unsupervised Pretraining

1. Introduction

Transfer Learning



It is generally not a good idea to train a very large DNN from scratch.

Transfer Learning

: Reusing the lower layers of an existing neural network

If the input picture don't have the same size as the ones used in the original task, you will have to resize them to the size expected by the original model.

2. Freezing the Lower Layers

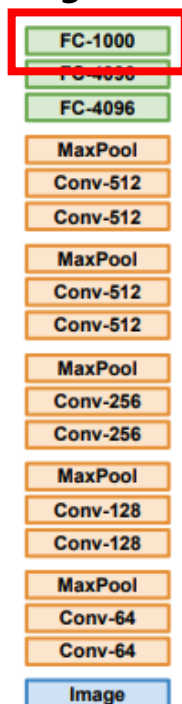
Small Dataset

Existing ConvNet



Reuse

Reusing ConvNet



Reinitialize & Train

- set for your problem.
(If you have C classes, then do FC-C.)
- replace and retrain the classifier

Frozen Layers

- freeze these weights when training the network.
- learning rate = 0
- during training, give the optimizer the list of variables to train, excluding the variables from lower layers.

2. Freezing the Lower Layers

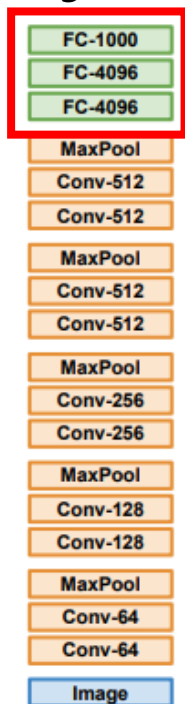
Bigger Dataset

Existing ConvNet



Reuse

Reusing ConvNet



Reinitialize & Train

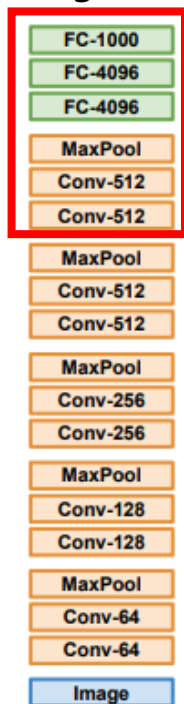
- fine-tune the weights of some higher level portion of pretrained network.

Frozen Layers

3. Tweaking, Dropping, or Replacing the Upper Layers

Tweaking

Reusing ConvNet



Reinitialize & Train

- the observation that the earlier features of a ConvNet contain more generic features (e.g. edge detectors or color blob detectors) that should be useful to many tasks, but later layers of the ConvNet becomes progressively more specific to the details of the classes contained in the original dataset.
- lower learning rate when fine tuning.
(1/10 of original learning rate is good starting point.
Especially, On Conv, more lower learning rate)

Frozen Layers

3. Tweaking, Dropping, or Replacing the Upper Layers

Tweaking

ConvNet



Specific



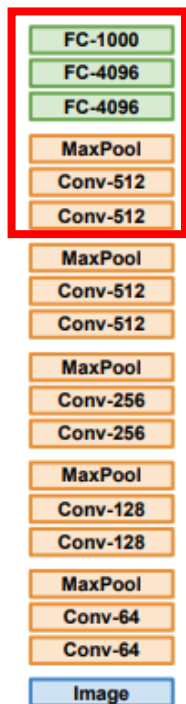
Generic

	Very similar dataset	Very different dataset
Very little data	Use Linear Classifier on top layer	You're in trouble... Try linear classifier from different stages
Quite a lot of data	Finetune a few layers	Finetune a larger number of layers

3. Tweaking, Dropping, or Replacing the Upper Layers

Dropping & Replacing

ConvNet



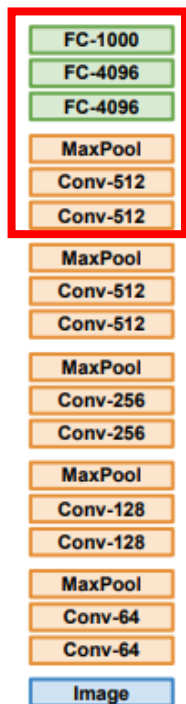
Reinitialize & Train with dropping or replacing

- If you cannot get good performance, and you have little training data, try dropping the top hidden layers.
- If you have plenty of training data, you may try replacing the top hidden layers instead of dropping them, and even add more hidden layers

4. Caching the Frozen Layers

Caching

ConvNet



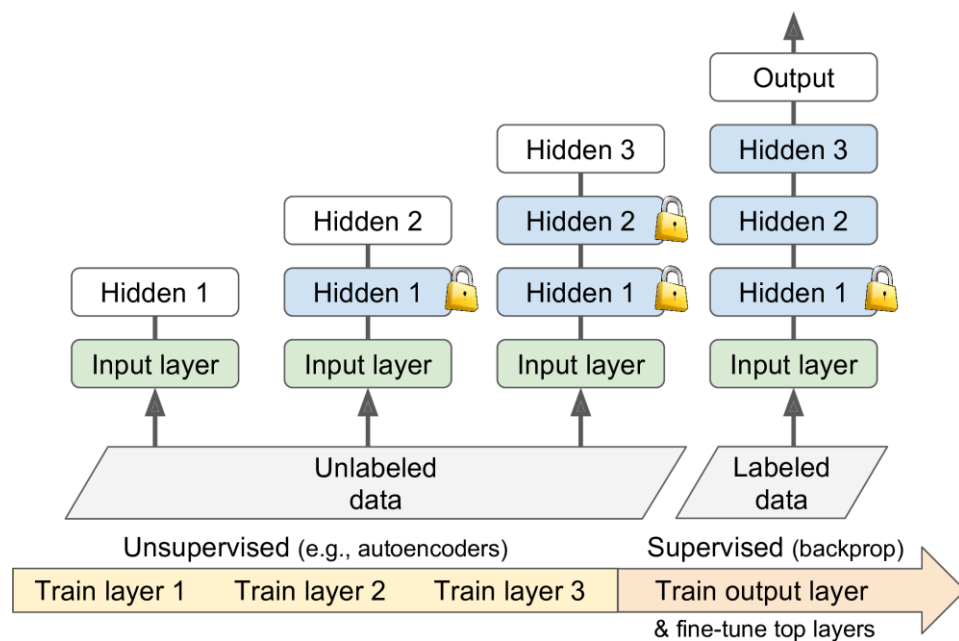
Reinitialize & Train with Caching

- ① run the whole training set through the lower layers
 - ② during training, instead of building batches of training instances, you would build batches of outputs from hidden layers and feed them to the training operation.
- It will give you a huge speed boost

Frozen Layers

5. Unsupervised Pretraining

Plenty of unlabeled data



If you don't have much labeled training data, you may be able to perform unsupervised pretraining.

- ① try to train the layers one by one, starting with the lowest layer and then going up, using an unsupervised feature detector (e.g. RBMs, autoencoders)
- ② Once all layers have been trained, you can finetune the network using supervised learning